

The Reliability Gap in LLM-as-a-Judge Evaluation: Agreement, Bias, and the Push for Calibrated Uncertainty (2024-2026)

Introduction

LLM-as-a-Judge (LLMaaJ) has become the default evaluation paradigm for large language model outputs, replacing or supplementing human review in scenarios where ground truth is unavailable or human annotation cannot scale ([Confident AI, 2026](#)). The appeal is straightforward: no annotation team, no weeks-long turnaround, just a prompt and an API call ([Confident AI, 2026](#)). But a growing body of research through mid-2026 shows that the metrics used to validate these judges are themselves flawed, and that the judges exhibit systematic, reproducible biases that exact-match agreement numbers conceal. The evidence points to a field that has scaled evaluation faster than it has validated it.

How Reliable Is LLM-as-a-Judge, Really?

Reported agreement numbers vary widely depending on task and methodology, and that variance is itself informative. On general-purpose evaluation prompts, GPT-4-style judges have shown up to roughly 85% agreement with human annotators, a figure often cited as exceeding the ~81% agreement rate between two human annotators on the same tasks ([Confident AI, 2026](#); [Vongthongsri, 2026](#)). On expert knowledge tasks, however, subject matter experts agreed with LLM judges only 64-68% of the time, meaning the judge was simply wrong by expert standards roughly one time in three ([Ng, n.d.](#)). A hands-on tutorial from Evidently AI found that an initial, unrefined judge prompt matched human expert labels only 67% of the time before iteration ([Evidently AI, 2025](#)).

This spread of 64% to 85% is not simply noise. It reflects the fact that the difficulty of the underlying task, and how "agreement" is measured, drives the number as much as judge quality does. The largest and most methodologically rigorous study to date makes this explicit.

The Kappa Deflation Problem

Norman, Rivera, and Hughes (2026) conducted what is currently the largest systematic evaluation of LLM-as-a-Judge: 21 judges from nine providers, evaluated across MT-Bench, JudgeBench, and RewardBench under three separate protocols (agreement, consistency, and bias audit), covering 118 runs and roughly 541,000 individual judgments ([Norman et al., 2026](#)). Their central finding is that the field's standard validation metric, exact-match agreement, systematically overstates a judge's discriminative ability because it does not correct for chance agreement. When the same judgments are re-scored using Cohen's kappa, a chance-corrected agreement statistic, scores drop by 33 to 41 percentage points on MT-Bench alone ([Norman et al., 2026](#)). This "kappa deflation" was universal across the full cohort of judges tested, including frontier models evaluated as recently as April 2026, indicating the problem is not confined to weaker or older systems ([Norman et al., 2026](#)).

In practical terms: a judge reported as "85% accurate" using exact match may, once chance agreement is removed, be closer to 45-50% in genuine discriminative skill. This is a substantial reframing of how much trust the field's most commonly cited agreement figures deserve.

The same study also found that judge rankings are not stable across benchmarks, shifting by up to 14 positions depending on which benchmark (MT-Bench, JudgeBench, or RewardBench) is used to rank them ([Norman et al., 2026](#)). A judge that ranks near the top on one benchmark can rank near the bottom on another. This undermines the common practice of selecting a single "best" judge model based on one leaderboard and deploying it broadly.

Perhaps most concerning, Norman et al. (2026) identified two production-deployed judges that combine high test-retest reliability (correlation above 0.95, meaning the judge gives consistent answers to the same input over repeated runs) with severe position bias (greater than 0.10) ([Norman et al., 2026](#)). This is a critical distinction: **reliability is not validity**. A judge can be highly self-consistent while consistently wrong in a systematic direction. High

test-retest scores can create false confidence that a judge is trustworthy, when in fact it is reliably biased rather than reliably correct.

Systematic Biases in LLM Judges

Independent of the agreement-metric problem, LLM judges exhibit several well-documented, recurring biases that distort scores regardless of how agreement is measured.

Position bias. Judges tend to favor whichever response is presented first (or, depending on the model, second) in a pairwise comparison, independent of actual quality. A systematic study published at IJCNLP 2025 found that the choice of judge model has a larger effect on positional bias than task complexity, output length, or the actual quality gap between responses, and that swapping response positions caused GPT-4's judgment to flip entirely in some cases ([Galileo AI, 2026](#)). This aligns with Norman et al.'s (2026) finding of severe position bias (>0.10) in production-deployed judges ([Norman et al., 2026](#)).

Self-enhancement (self-preference) bias. Research presented at NeurIPS 2024 found that LLM evaluators recognize and systematically favor outputs generated by themselves or closely related models, with a proven linear correlation between a model's self-recognition capability and the strength of its self-preference bias ([Galileo AI, 2026](#)). This has direct implications for any pipeline where a model family is used to both generate and evaluate outputs.

Verbosity bias. Longer responses tend to receive higher scores from LLM judges even when shorter responses are more accurate and more useful ([Ng, n.d.](#)).

The evaluation-generation asymmetry. A 2024 paper from Seoul National University, dubbed "The Generative AI Paradox," found that LLMs perform significantly worse on evaluation tasks than on generation tasks: a model can often produce a strong answer but cannot reliably judge whether an answer (its own or another's) is actually good ([Ng, n.d.](#)).

Summary of Reliability Findings

Study / Source	Metric	Finding
Norman et al. (2026)	Exact match vs. Cohen's kappa	33-41 pp deflation on MT-Bench; universal across 21 judges (link)
Norman et al. (2026)	Cross-benchmark ranking stability	Rankings shift by up to 14 positions (link)
Norman et al. (2026)	Test-retest reliability vs. position bias	>0.95 reliability coexists with >0.10 position bias in 2 production judges (link)
Confident AI (2026)	Human-label agreement	~85% agreement, vs. ~81% human-human agreement (link)
Expert knowledge tasks (ACM IUI 2025, via Ng)	Human-label agreement	64-68% agreement with subject matter experts (link)
Galileo AI (2026)	Implementation survey	93% of teams report struggling with LLM judge implementation (link)

Approaches to Quantifying Judge Uncertainty

Because exact-match agreement is now understood to be an unreliable validation signal, researchers have developed methods to directly quantify the uncertainty of individual judge outputs rather than relying solely on aggregate agreement statistics.

Black-box confusion-matrix methods. Wagner et al. (2024) introduced a black-box uncertainty quantification method for LLM-as-a-Judge that requires no access to model internals. It analyzes the relationships between a judge's generated assessments and the space of possible ratings, cross-evaluates these relationships, and constructs a confusion matrix based on token

probabilities to derive high- or low-uncertainty labels for each judgment ([Wagner et al., 2024](#)). Their evaluation across multiple benchmarks found a strong correlation between the accuracy of LLM evaluations and the derived uncertainty scores, suggesting the method can flag unreliable judgments before they are trusted downstream ([Wagner et al., 2024](#)).

Linear probes on hidden states. A more recent and computationally efficient approach comes from Radharapu et al. (2025), who trained linear probes on reasoning judges' hidden states using a Brier score-based loss (a proper scoring rule for probabilistic calibration) to produce calibrated uncertainty estimates without any additional model training ([Radharapu et al., 2025](#)). Tested on both objective tasks (reasoning, mathematics, factuality, coding) and subjective human-preference judgments, the probes achieved better calibration than existing verbalized-confidence and multi-generation methods while using approximately 10x fewer computational resources ([Radharapu et al., 2025](#)). The tradeoff is that probes produce conservative estimates that underperform on easier datasets, though the authors note this conservatism may actually be desirable in safety-critical deployments that prioritize low false-positive rates over raw accuracy ([Radharapu et al., 2025](#)). This method generalized robustly to evaluation domains it was not trained on, which addresses a common failure mode where uncertainty methods work only in-distribution ([Radharapu et al., 2025](#)).

Conformal prediction. Broader LLM uncertainty research, some of it directly applicable to judge reliability, has moved toward conformal prediction (CP) methods that provide statistically guaranteed uncertainty bounds rather than raw confidence scores. Selective Conformal Uncertainty (SConU), published at ACL 2025, uses conformal p-values to determine whether a given output falls within a model's calibrated uncertainty distribution, while Adaptive Conformal Prediction for Factuality (arXiv:2604.13991) extends this with prompt-dependent calibration via conditional quantile regression rather than relying on marginal coverage guarantees alone ([Zylos Research, 2026](#)). The practical appeal for production systems is interpretability: a system can report "95% confident the answer is in this set" with an empirically-verified guarantee, rather than an unvalidated confidence number ([Zylos Research, 2026](#)).

Approaches to Improving Judge Calibration and Alignment

Beyond quantifying uncertainty after the fact, the field has developed several practical strategies to improve judge reliability upstream.

Decomposing scores into binary verdicts. Rather than asking a judge for a numeric score on a scale (e.g., 1-10), which is highly sensitive to prompt phrasing and model-specific scale interpretation, researchers recommend decomposed binary yes/no questions such as "does this response contain a hallucinated fact?" Galileo AI's ChainPoll technique, which combines binary verdicts with repeated polling, reportedly improved accuracy by 23% over single-shot numeric scoring ([Ng, n.d.](#)). Evidently AI similarly recommends binary or few-class categorical labels over continuous scores as a default practice ([Evidently AI, 2025](#)).

Pairwise preference alignment. Liu et al.'s Pairwise Preference Search (PairS) systematically analyzes the inconsistencies and biases LLM evaluators exhibit and proposes a strategy to more closely align pairwise judgments with human perspectives, addressing limitations in prior calibration approaches ([EmergentMind summary of Liu et al., n.d.](#)).

Human-in-the-loop calibration. Multiple sources converge on the same operational recommendation: LLM judges should be treated as a way to scale human judgment, not replace it. The recommended workflow is to first label a representative sample of outputs manually to define what "quality" means for a specific use case, then iteratively refine the judge prompt against those human labels, measuring alignment with standard classification metrics (accuracy, precision, recall) until performance is acceptable ([Evidently AI, 2025](#)). When judge and human evaluations diverge, the recommendation is to trust the human and retune the judge, not the reverse ([Ng, n.d.](#)).

Multi-judge consensus. Rather than relying on a single judge model, some production teams use ensembles of multiple judges and require consensus, a practice reported to achieve 97-98% accuracy in elite implementations, compared to substantially lower rates for single-judge setups ([Galileo AI, 2026](#)).

Decomposed, dimension-specific evaluation. Splitting an overall "quality" assessment into concrete, independently measurable dimensions (factual accuracy, relevance, completeness, tone) allows some dimensions to be checked deterministically instead of relying on LLM judgment for everything, reducing the surface area where judge unreliability can enter the pipeline ([Ng, n.d.](#); [Evidently AI, 2025](#)).

Broader Calibration Context: Why This Is Hard

The reliability problems in LLM-as-a-Judge are a specific instance of a broader calibration deficit across LLMs generally. Alignment training, particularly RLHF, systematically degrades calibration by rewarding confident-sounding answers regardless of whether the model actually possesses the underlying knowledge ([Zylos Research, 2026](#)). As LLMs are increasingly deployed as autonomous agents taking longer action chains, a miscalibrated confidence signal compounds into a root cause of cascading failures, rather than a minor inconvenience ([Zylos Research, 2026](#)). Zylos Research (2026) frames calibration as a product of the entire training pipeline rather than something to be patched at inference time, recommending that teams doing domain fine-tuning add calibration measurement directly to their evaluation suite rather than discovering miscalibration only after deployment ([Zylos Research, 2026](#)).

There is also a feedback-loop risk specific to using LLM judges to drive further model training: if judge scores are used as a fine-tuning signal, models learn to optimize for what pleases the judge (verbosity, hedging, stylistic patterns the judge rewards) rather than genuine output quality, a direct instance of Goodhart's Law ([Ng, n.d.](#)).

Analysis and Conclusion

Taken together, the 2024-2026 literature supports a fairly clear and somewhat uncomfortable conclusion: **LLM-as-a-judge is useful at scale but has not earned the level of trust the field has extended to it**, largely because the field has validated it with the wrong metric. Exact-match agreement, still the dominant reporting standard, systematically overstates discriminative power by tens of percentage points once chance agreement is properly removed ([Norman et al., 2026](#)). This is not a minor statistical footnote, it changes the

interpretation of nearly every previously reported "judge accuracy" figure in the literature. Combined with evidence that judge rankings are benchmark-dependent and unstable, and that high self-consistency can mask severe, systematic bias rather than indicate trustworthiness, the practical implication is that no single judge model or single benchmark score should be treated as authoritative ([Norman et al., 2026](#)).

The most credible path forward combines three elements that appear consistently across the reviewed sources: chance-corrected, multi-benchmark validation of judges rather than single exact-match figures; direct uncertainty quantification at the level of individual judgments, via black-box confusion-matrix methods or lightweight hidden-state probes, so that low-confidence judgments can be flagged or routed to human review; and continued human-in-the-loop calibration rather than one-time validation, since judge alignment appears to drift and vary by domain ([Wagner et al., 2024](#); [Radharapu et al., 2025](#); [Evidently AI, 2025](#)). Multi-judge consensus and binary/decomposed verdicts are the most practically validated mitigations currently in use, with reported gains large enough (23% accuracy improvement for ChainPoll-style polling, 97-98% accuracy for consensus ensembles) to justify their added cost and latency in high-stakes deployments ([Ng, n.d.](#); [Galileo AI, 2026](#)). Given that 93% of teams report struggling with LLM judge implementation despite near-universal adoption, the gap between how judges are actually used in production and how rigorously they have been validated remains the central unresolved problem in this area as of mid-2026 ([Galileo AI, 2026](#)).

References

Confident AI. (2026, May 16). LLM-as-a-judge simply explained: The complete guide to run LLM evals at scale. <https://www.confident-ai.com/blog/why-llm-as-a-judge-is-the-best-llm-evaluation-method>

Evidently AI. (2025, November 14). How to align LLM judge with human labels: A hands-on tutorial. <https://www.evidentlyai.com/blog/how-to-align-llm-judge-with-human-labels>

Galileo AI. (2026). LLM-as-a-judge vs human evaluation. <https://galileo.ai/blog/llm-as-a-judge-vs-human-evaluation>

Liu, et al. (n.d.). Pairwise preference for LLM evaluator alignment [Summary]. EmergentMind. <https://www.emergentmind.com/papers/2403.16950>

Ng, J. W. (n.d.). The irony of testing AI. <https://jiaweing.com/blog/the-irony-of-testing-ai>

Norman, J. D., Rivera, M. U., & Hughes, D. A. (2026, June 17). Reliability without validity: A systematic, large-scale evaluation of LLM-as-a-judge models across agreement, consistency, and bias. arXiv. <https://arxiv.org/html/2606.19544v1>

Radharapu, B., Saxena, E., Li, K., Whitehouse, C., Williams, A., & Cancedda, N. (2025, December 23). Calibrating LLM judges: Linear probes for fast and reliable uncertainty estimation. arXiv. <https://arxiv.org/html/2512.22245>

Wagner, N., Desmond, M., Nair, R., Ashktorab, Z., Daly, E. M., Pan, Q., Santillán Cooper, M., Johnson, J. M., & Geyer, W. (2024, October 15). Black-box uncertainty quantification method for LLM-as-a-judge. arXiv. <https://arxiv.org/pdf/2410.11594v1>

Zylos Research. (2026, April 18). LLM calibration and uncertainty quantification in production AI agents. <https://zylos.ai/research/2026-04-18-llm-calibration-uncertainty-production-agents/>